

LoRa Reverse Engineering and AES EM Side-Channel Attacks using SDR

Pieter Robyns

About me

- PhD student at Hasselt University since 2014
 - Since 2016 on FWO SBO research grant
- Researching wireless security
 - Protocol security, location tracking, fingerprinting
 - Machine learning and side channel analysis
 - Wi-Fi, GSM, LoRa, proprietary protocols
- Website: <https://robyns.me>

Email: pieter.robyns@uhasselt.be



Motivation for researching LoRa

- Project started in April 2016 → LoRa was relatively new
 - Introduced to LoRa by co-advisor
- A lot of opportunities to learn new things
 - No working software-based decoders available, only simulations
 - Building a GNU Radio OOT module from scratch
 - Limited description of the PHY layer: patents and blog posts
 - Reverse engineering low-level aspects of a protocol
 - Fingerprinting and tracking devices over long ranges
 - Machine learning applied to fingerprinting instead of expert feature selection
 - Side-channel attacks
 - IoT devices are inherently more vulnerable



Part 1

Unlocking the LoRa PHY

Unlocking the LoRa PHY

- Hardware LoRa radios can only be interfaced with over a serial connection

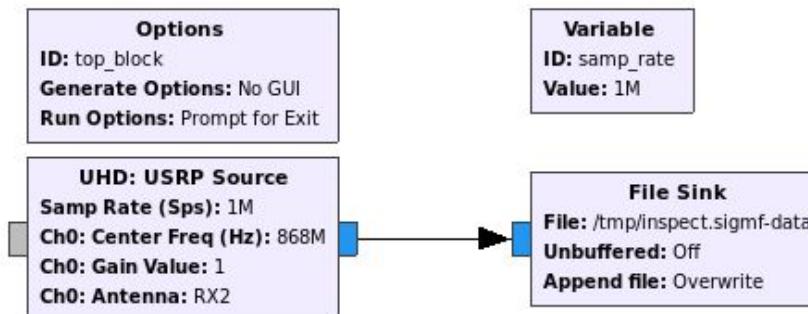


Microchip RN2483 + custom
board made by my co-advisor

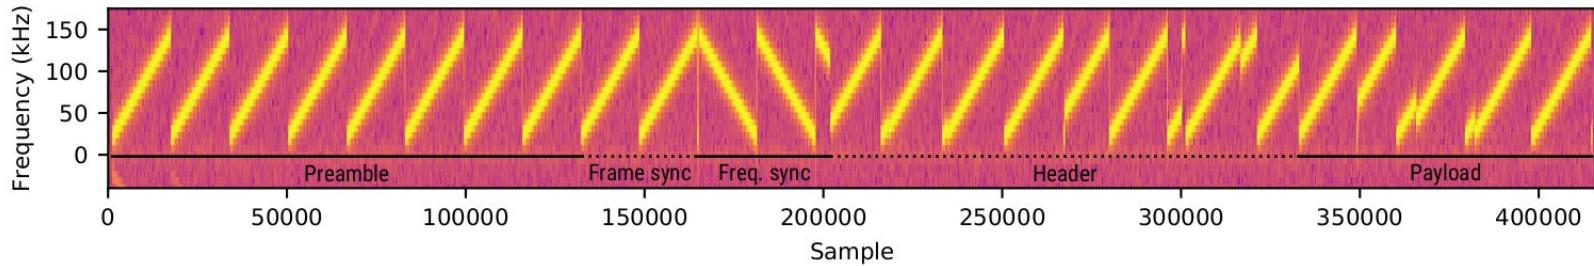
- We need access to the raw PHY signal for fingerprinting
⇒ Where do we start?

Unlocking the LoRa PHY

- GNU Radio to the rescue! Let's inspect a transmission using a simple flowgraph



Unlocking the LoRa PHY



- Frame structure can be easily derived from patent
 - See [Patent EP2763321 A1](#)
 - Also contains information on:
 - Modulation
 - Interleaving
 - Some other info located in datasheets:
 - Whitening and coding
- Let's build a receiver!

Patents

Find prior art Discuss this application Vie

Low power long range transmitter
EP 2763321 A1

ABSTRACT

A transmitter device arranged to encode a set of digital input data into a succession of modulated chirps, whereby said digital input data are encoded according to a Gray code into codewords (320, 321, 322) having a plurality of bits, and having an interleaver that distributes the bits (C_{00}, \dots, C_{M1}) of each codeword into a series of digital modulation values (S_0, \dots, S_7), at different bit positions, and to synthesize a series of modulated chirps whose cyclical shifts are determined by the modulation values. A special frame structure is defined in order to ensure high robustness, and variable bit-rate flexibility.

Publication number EP2763321 A1
Publication type Application
Application number EP20130154071
Publication date Aug 6, 2014
Filing date Feb 5, 2013
Priority date Feb 5, 2013
Also published as CN103973626A, US9252834, US20140219329
Inventors Olivier Bernard André Seller, Nicolas Somin
Applicant Semtech Corporation
Export Citation BibTeX, EndNote, RefMan
Patent Citations (7), Referenced by (8), Classifications (9), Legal Events (5)
External Links: Espacenet, EP Register

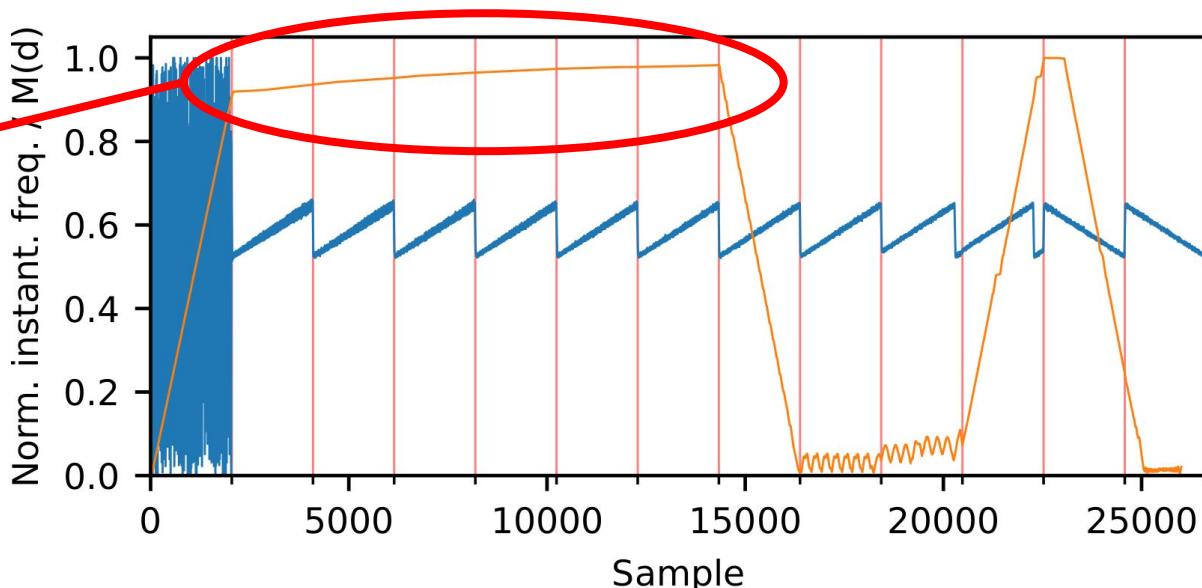


How do we detect the signal?

- Detecting: pretty standard problem in signal processing
- Multiple solutions possible; I chose Schmidl-Cox algorithm
 - Autocorrelation exploiting the repeating property of the preamble

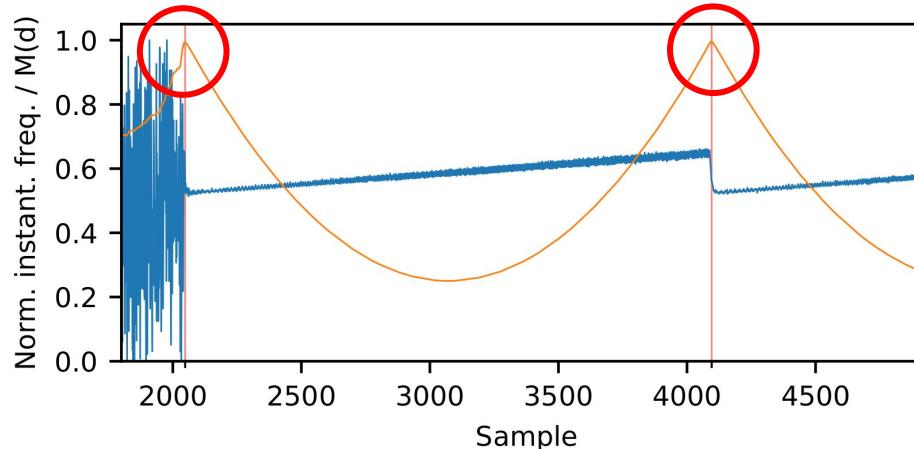
Preamble is here,
but where does it
start?

Thresholding = bad!



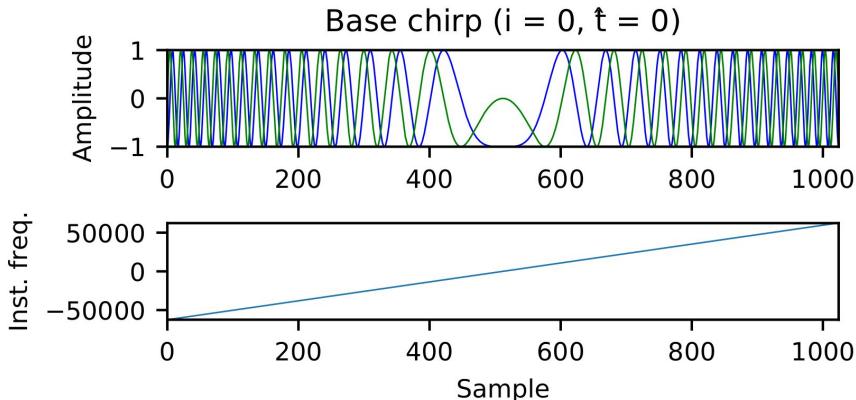
How do we synchronize to the signal?

- Again multiple possibilities:
 - Demodulate preamble symbol → supposed to be 0
 - Offset from 0 indicates a time shift (basic principle of LoRa modulation as we will see)
 - However: ambiguity because a frequency shift also causes an offset from 0!
 - Cross-correlate instantaneous frequency with locally generated preamble
 - Higher sensitivity to noise, but no ambiguity

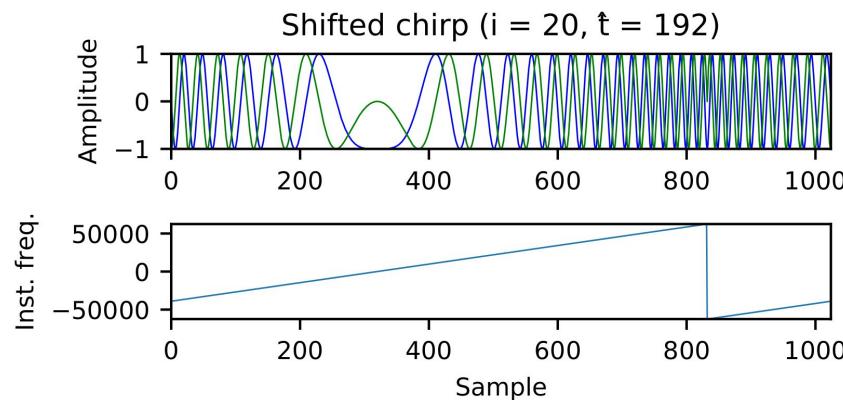


How do we demodulate a single symbol?

- Modulation of LoRa is based on Chirp Spread Spectrum
- Chirp = signal that linearly increases in frequency
- To modulate a value “i” onto chirp: cyclically time shift it!



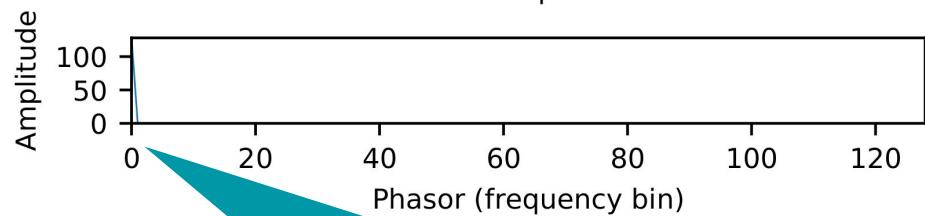
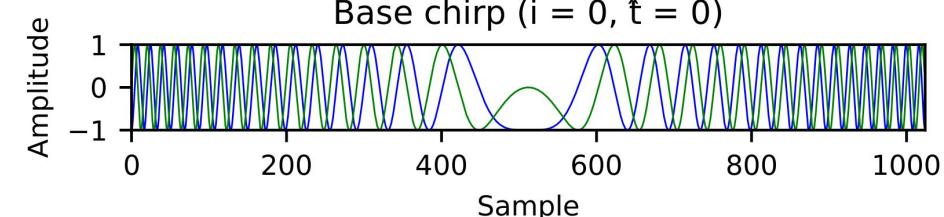
Value: 0 (unmodulated)



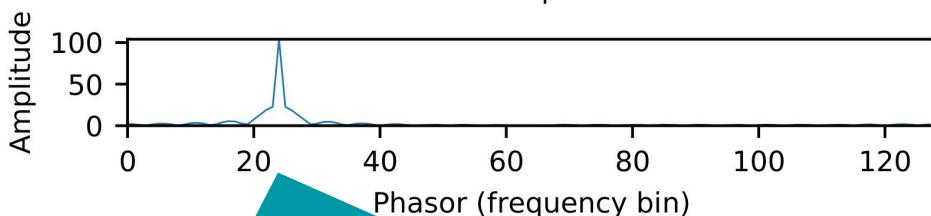
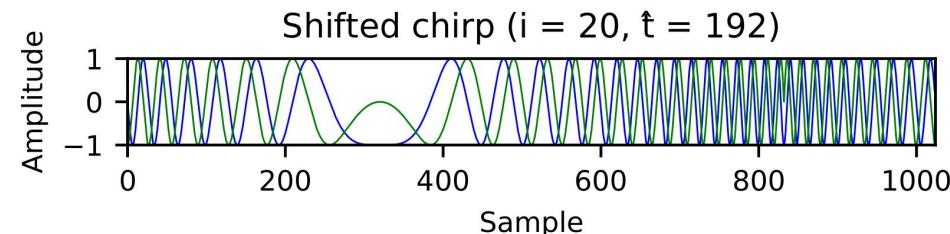
Value: 20 (spoiler: indexing ;))

How do we demodulate a single symbol?

- Cyclic shift results in a peak in the frequency domain when multiplied by a conjugate base chirp (+ resampling at chirp rate) \Rightarrow details not important for now
- Index is “gray” decoded. Encode to demodulate!



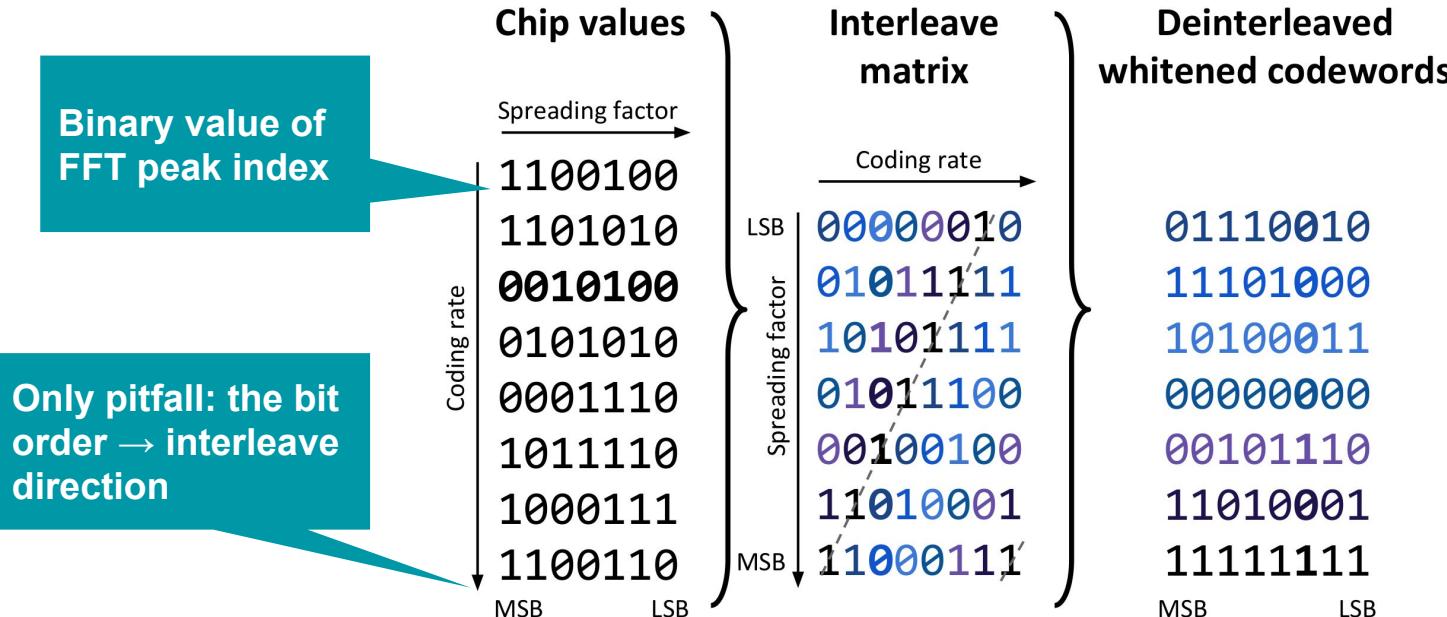
gray(0) == 0 == i



gray(24) == 20 == i

Demodulation continued: interleaving

- Interleaving is trivial: algorithm provided in patent
 - Spreading factor determines bits per symbol value (here: 7)
 - Coding rate determines symbol values per interleave matrix (here: 8)

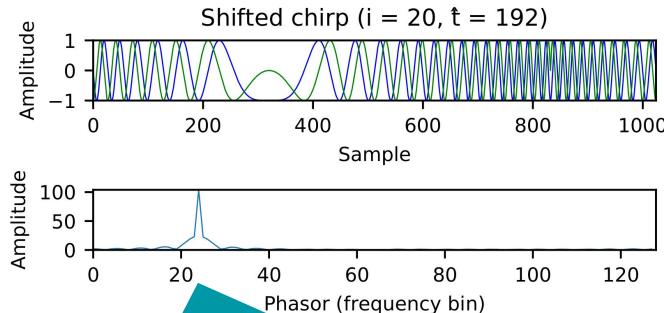


Unlocking the LoRa PHY: unknown aspects

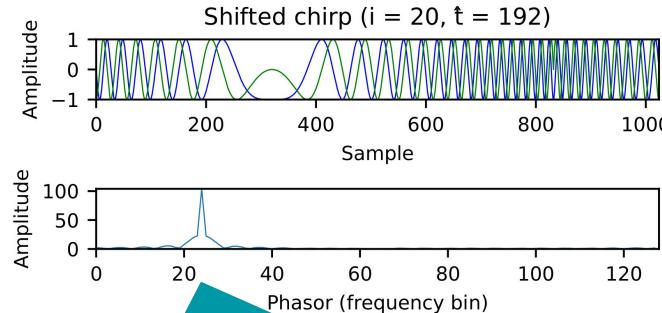
- What's left to be done?
 - ~~How do we detect the signal?~~
 - ~~How do we synchronize to the signal?~~
 - ~~How does the modulation and interleaving work?~~
 - What is the relation between a raw symbol and its integer value?
 - In which stage of the decoding is whitening performed and how?
- Not discussed in this presentation:
 - Header structure
 - Clock drift correction
 - Swapping of nibbles + CRCs
 - See my paper for more info!

Relation between symbol and integer value?

- Patent states “gray coding” is used
 - Total of 4 possible mappings to symbol values:



gray(24) or degrey(24)?



gray(103) or degrey(103)?

- To check correctness: implement decoder up to interleaving and look for patterns
 - Header is unwhitened \Rightarrow use header to check previous stages

c. Relation between symbol and integer value?

- Example: sending packets with increasing payload sizes (SF 7)

	Bin data	Gray encoding	Gray decoding
Hex len			
Right to left			
(FFT bin)			
127 → 0			
11:	01: 10001100 00001000 10000011 01000010 00101000 02: 10001100 00001000 01100100 01000010 00101001 03: 10001100 00001000 00000111 01000010 00100000 10: 10001100 10000010 01100011 01000001 00100001 11: 10001100 10000010 00000000 01000001 00101000 12: 10001100 10000010 11100111 01000001 00101001 20: 10001100 10000010 10100110 00000000 00100001 21: 10001100 10000010 11000101 00000000 00101000 22: 10001100 10000010 00100010 00000000 00101001	00001100 01001010 10000011 01000010 00000000 10001000 01001010 01000101 01000010 00100001 10001000 01001010 00000111 01100010 00001000 10001100 11000010 01000010 01000001 00100001 10001100 10000010 00100000 01000001 00101000 00001000 11000010 11000110 01000001 00101001 00001000 10000010 10000111 00000000 00100001 00001000 10000010 11100101 00000000 00101000 10001100 10000010 00000011 00000000 00101001	
Left to right			
(FFT bin)			
0 → 127			
10:	01: 00000000 10001011 10011100 00000000 10001011 02: 00000000 01001110 10011100 00000000 00101101 03: 00000000 11000110 10011100 00000000 01001110 10: 10001011 00000000 10011100 10001011 11111111 11: 10001011 10001011 10011100 10001011 10011100 12: 10001011 01001110 10011100 10001011 01100011 20: 01001110 00000000 10011100 10001011 00111010 21: 01001110 10001011 10011100 10001011 01011001 22: 01001110 01001110 10011100 10001011 10100110	10011000 10001011 10011000 00100000 00011110 00011100 01001110 01111100 00100000 11100000 00011100 11000101 00011110 00000000 10001010 10010011 10001000 01111011 10011000 11110111 10010011 00000011 00011001 10011000 10011101 00010111 11000110 11111111 10011000 01100011 11010010 11001000 10111110 11011001 00110010 11010010 01000011 11011100 11011001 01011000 01010110 10000110 00111010 11011001 10100110	

How do we decode the obtained codewords?

```
01: 00000000 10001011 10011100 00000000 10001011  
02: 00000000 01001110 10011100 00000000 00101101  
03: 00000000 11000110 10011100 00000000 01001110  
10: 10001011 00000000 10011100 10001011 11111111  
11: 10001011 10001011 10011100 10001011 10011100  
12: 10001011 01001110 10011100 10001011 01100011  
20: 01001110 00000000 10011100 10001011 00111010  
21: 01001110 10001011 10011100 10001011 01011001  
22: 01001110 01001110 10011100 10001011 10100110
```

- Coding: 4/5 - 4/8 as options imply Hamming coding
- Payload whitening: XOR with random LFSR
 - Mentioned but specified algorithm doesn't work in practice :(.
 - In what stage is the data whitened?
 - Only payload is whitened → very useful!

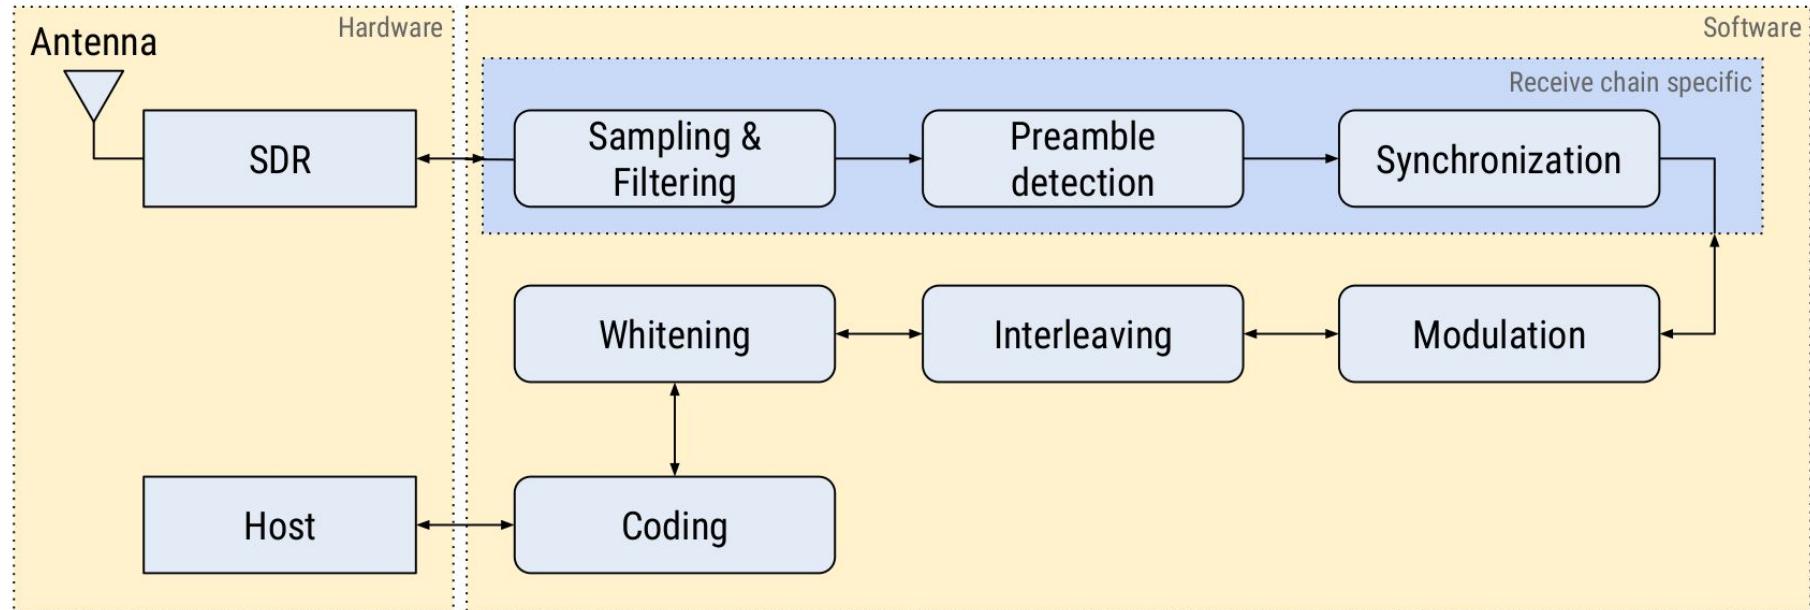
How do we decode the obtained codewords?

- Fastest solution: brute force
- Whitening: send payload with all zeros
 - Hamming code of 0000 is 00000000, which is convenient
 - Ideas for determining LFSR algebraically welcome!
- Hamming codes
 - Try all possible bit permutations for a header byte. Choose the one without decode errors
 - Verify with multiple (all possible) header byte values
 - $\frac{8!}{(8-4)!} = 1680$

10001011
↙

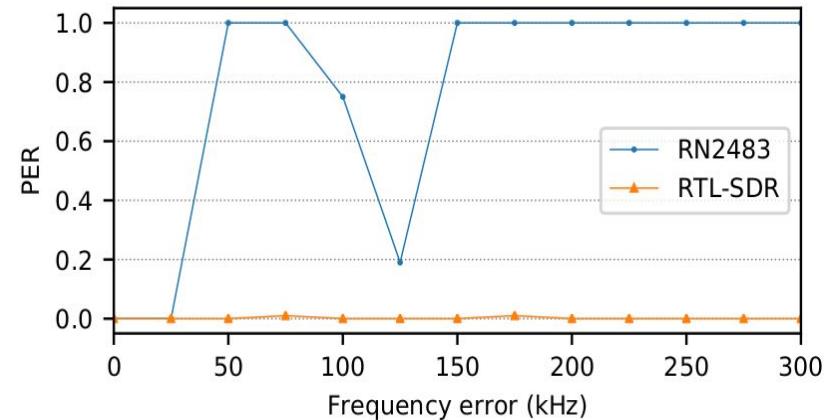
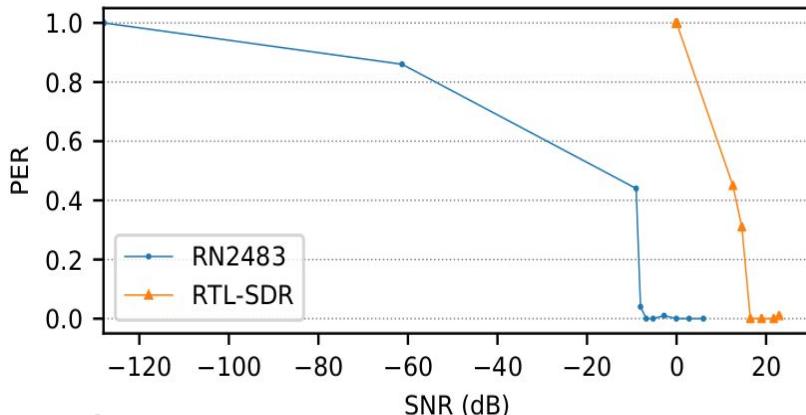
Results

- Overview of all components linked together:



Results

- Comparison with real hardware:



- Code: <https://github.com/rpp0/gr-lora>
 - Special thanks to my student William for implementing some optimizations
- Other decoders / related work
 - LoRa-SDR: <https://github.com/myriadrf/LoRa-SDR>
 - BastilleResearch's gr-lora: <https://github.com/BastilleResearch/gr-lora>

Application

Fingerprinting LoRa devices using neural networks

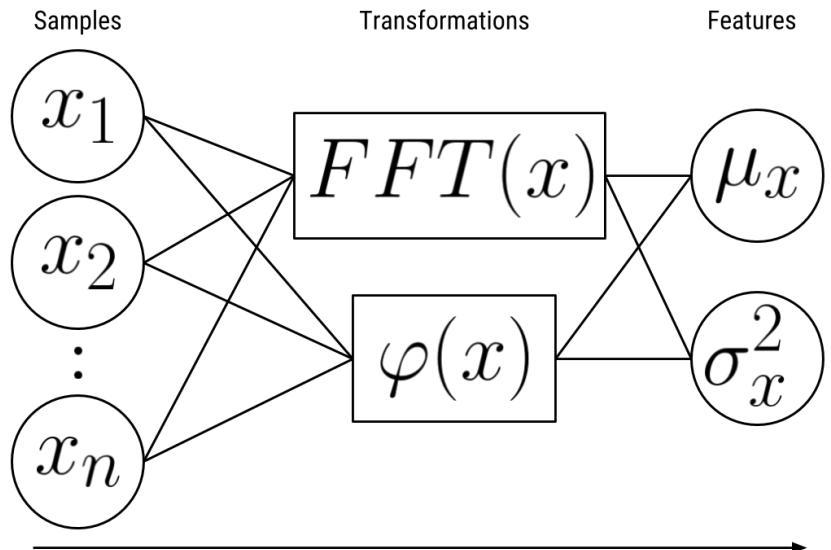
Why fingerprint devices?

- Defensive
 - Extra layer of defense in critical infrastructure → detect unknown devices
 - Possibly counter relay attacks
 - Measure degree of privacy provided by device
- Offensive
 - Linking anonymous transmissions (e.g. defeat MAC randomization)
 - Tracking the location of sensors (e.g. to take them down)
 - Mimic radio signature of a device to defeat IDSs
- Caveat: cat-and-mouse game between attacker and defender!

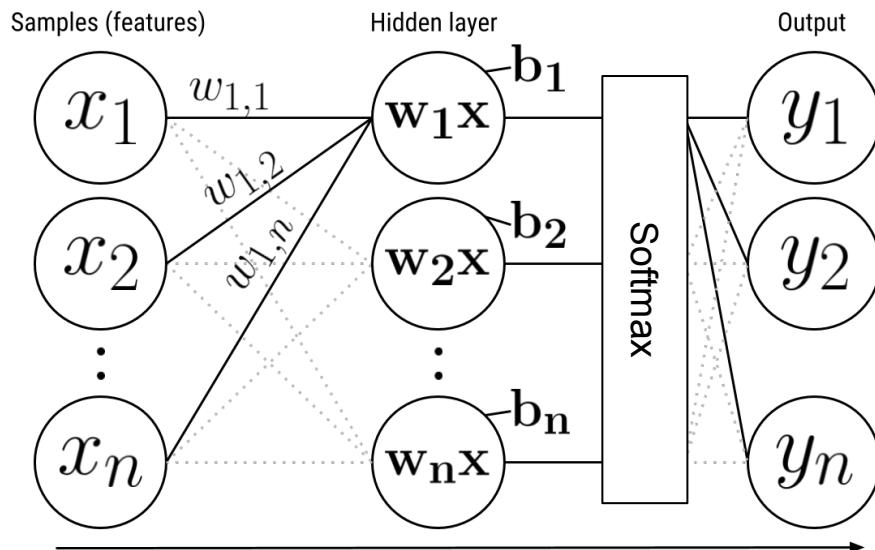
PHY-layer fingerprinting theory

- Hypothesis: no two radios can be perfectly identical
 - Manufacturing differences in circuits, crystal oscillators, components, ...
 - Manifest as per-device transmission errors (e.g. frequency offset)
 - Error tolerance typically defined within data sheets (e.g. ± 12 KHz)
 - *Larger tolerance implies more entropy*
- Challenge: distinguish noise from errors caused by the radio hardware
 - Traditional approach: use statistical measures on “expert features”
 - Carrier Frequency Offset, Sampling Frequency Offset, Preamble Transient,...
 - My approach: apply machine learning to the **raw radio signal**
 - Similar techniques applied in face recognition, image classification, etc.

Simplified comparison

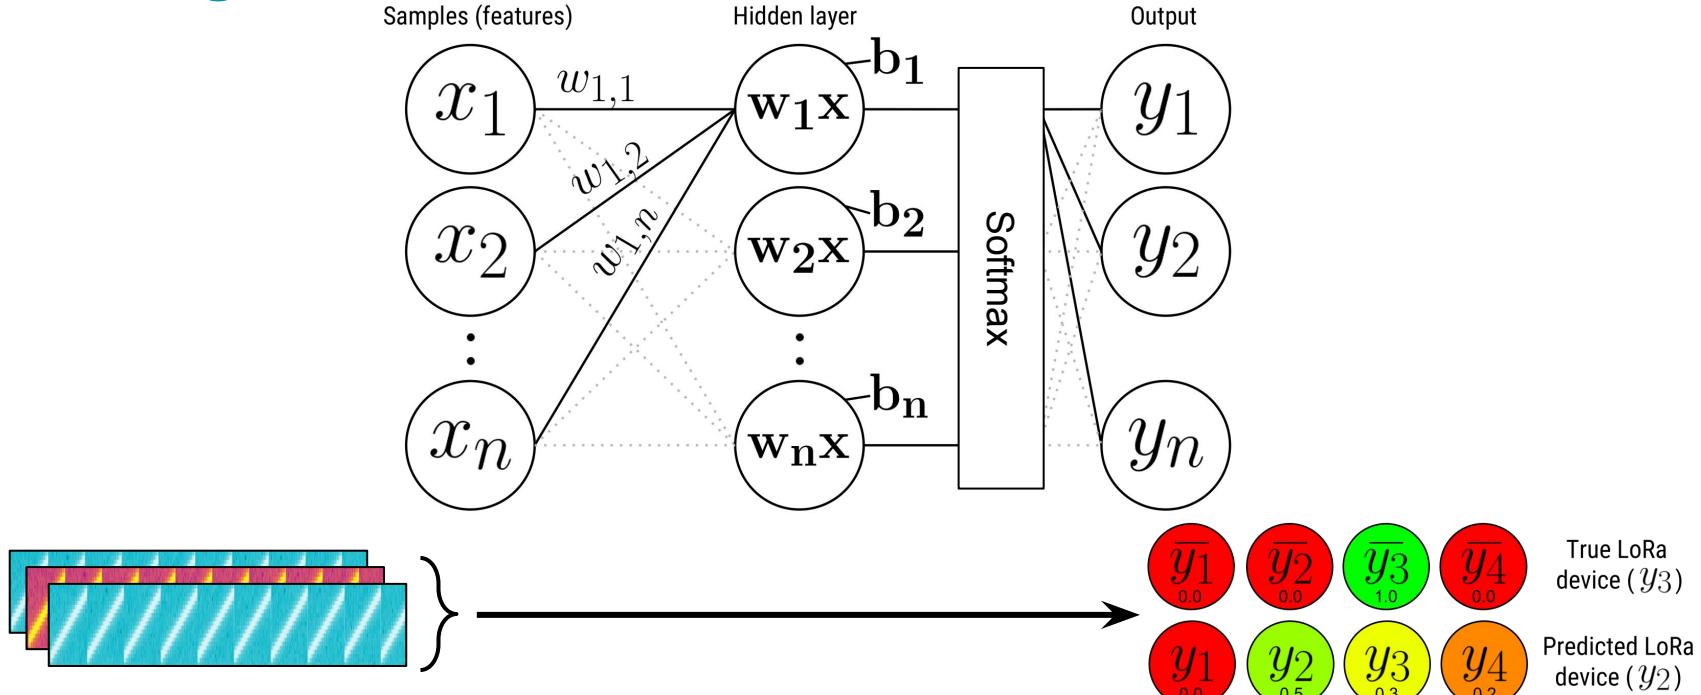


- “Human” filtering at feature level
- Resulting features can be learned with ML or statistical distance measures



- Unimportant features are filtered through weight values
- Consider raw samples as features

Training the neural network

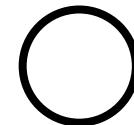
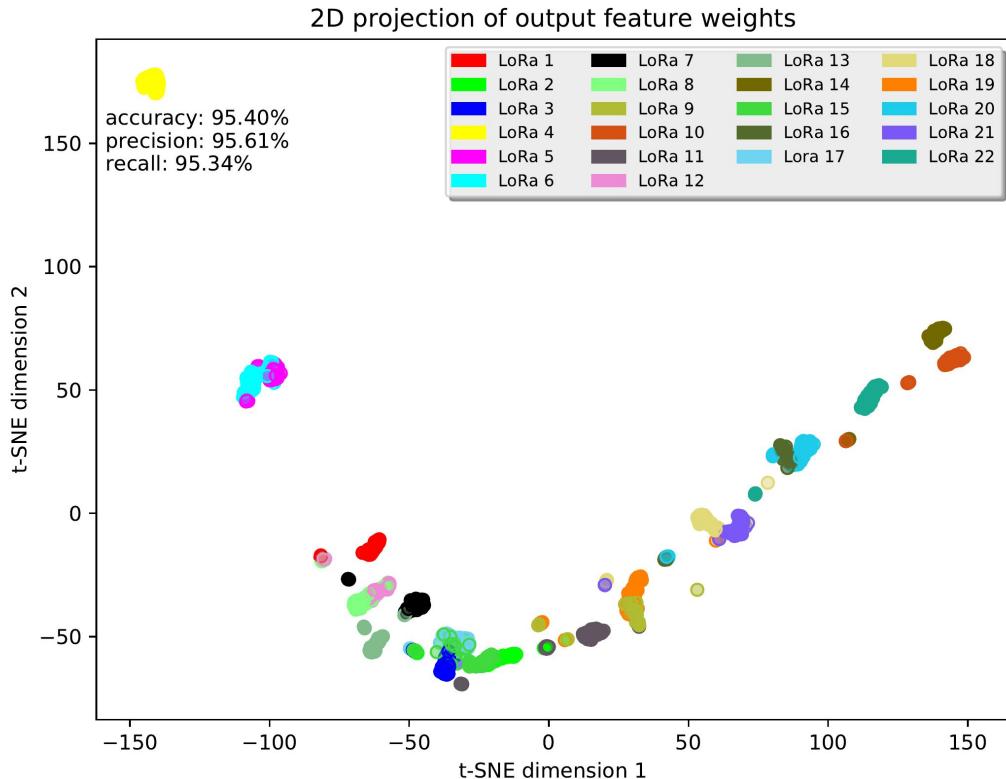


1. Label transmission with LoRa device.
2. Feed data through neurons and check resulting outputs.
3. Evaluate the result in terms of a “loss” function, and update the neuron weights accordingly. Repeat step 2.

LoRa fingerprinting experiment

- Experiment: can we uniquely identify 22 LoRa devices?
 - 3 different vendors
 - 1 SX1272
 - 2 RF96
 - 19 RN2483
 - Model: simple MLP from previous slides
 - Training data: ~100,000 symbols
 - Test data: ~1,000 symbols
- 95% accuracy
 - However: tradeoff between sensitivity to noise and being able to detect fine-grained differences between devices → noise is a problem

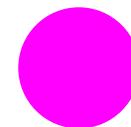
Results



Outline: predicted device



Fill: true device



Correct



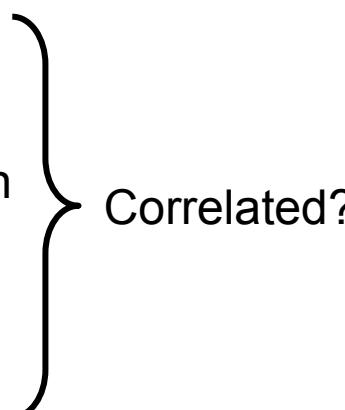
Incorrect

Each point is one symbol!
(>16 symbols per frame)

Part 2

EM side-channel attacks on AES

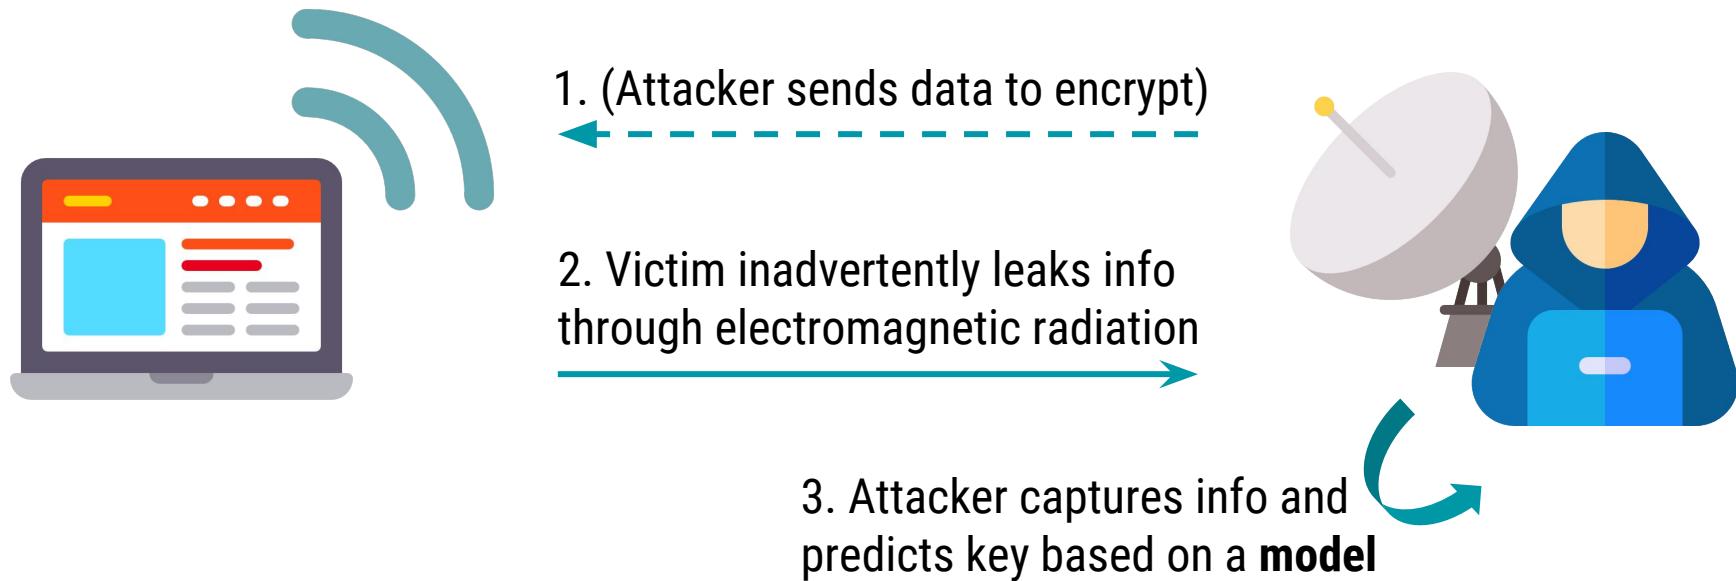
What is a side channel attack?

- Implementation leaks information through “side channel”
 - Attacker gains *advantage* based on this information
 - Numerous types of side channels:
 - Timing
 - Acoustic
 - Power consumption
 - Temperature
 - Cache
 - Electromagnetic
- 
- Correlated?

Motivation

- EM side-channel attacks (on AES) are interesting
 - Used by LoRa, Wi-Fi, TLS, IPsec, apps, ...
- Attack techniques have been around for quite some time, but expensive equipment often required
- Can we do these TEMPEST-style attacks with cheap SDRs?
 - We will discuss a simple Correlation Power Attack (more complicated attacks exist)

Examples of EM side channel attacks



EM models

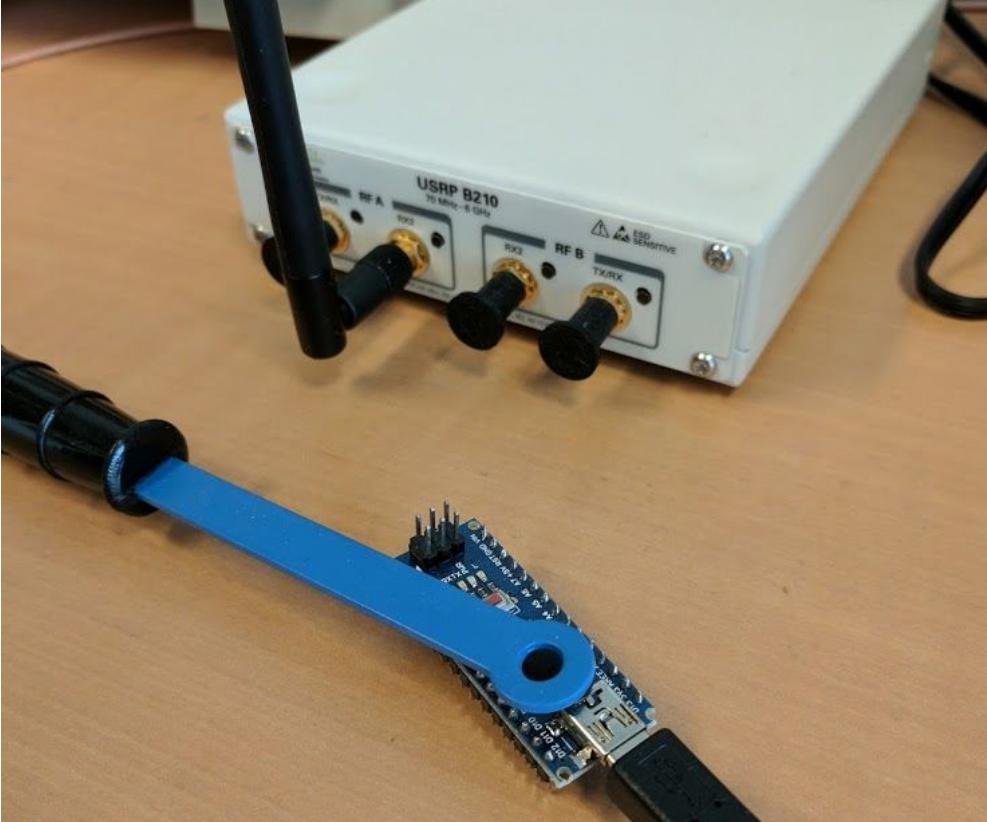
- Behavior of system can be approximated with a model
- Accuracy of model is crucial for successful attack
- Some observations:
 - Amplitude of electromagnetic radiation is proportional to power
 - Power is required to change state of a circuit

⇒ State changes cause variations in the amplitude of EM radiation, proportional to their power consumption
- What happens if we would AM demodulate AES encryptions?

Case: AES on ATmega 328p

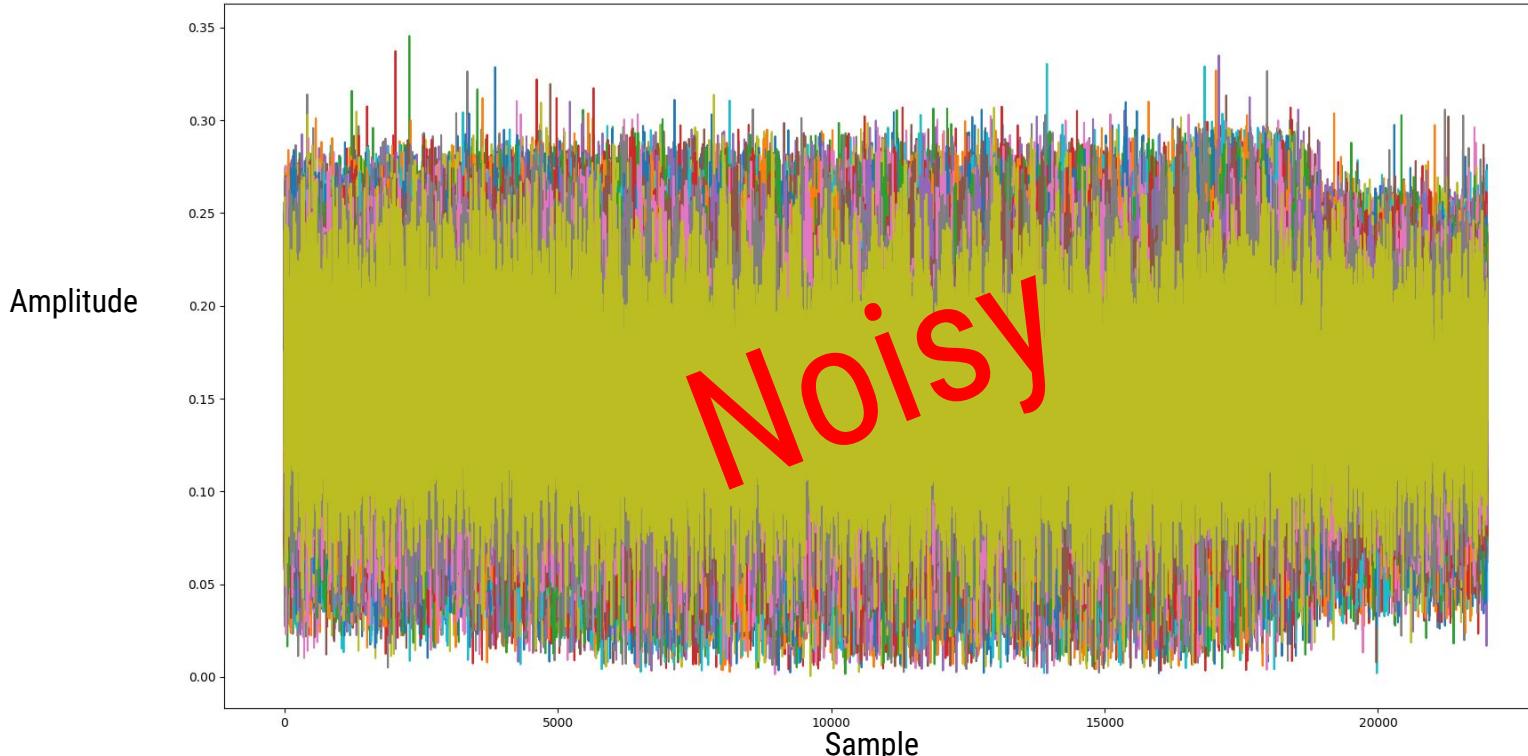
- Case study: AM demodulated AES encryptions performed by an ATmega 328p (Riscure competition)
 - Key size and key unknown; black box
- What we can learn from related works:
 - Lower frequencies must be favored^[1]
 - Harmonics of CPU clock frequency contain useful information^[2]
- Equipment: USRP B210 + amplifier + EM probe
 - ~18,000 traces. More = better

Case: AES on ATmega 328p



Case: AES on ATmega 328p

- AM demodulation of raw capture:



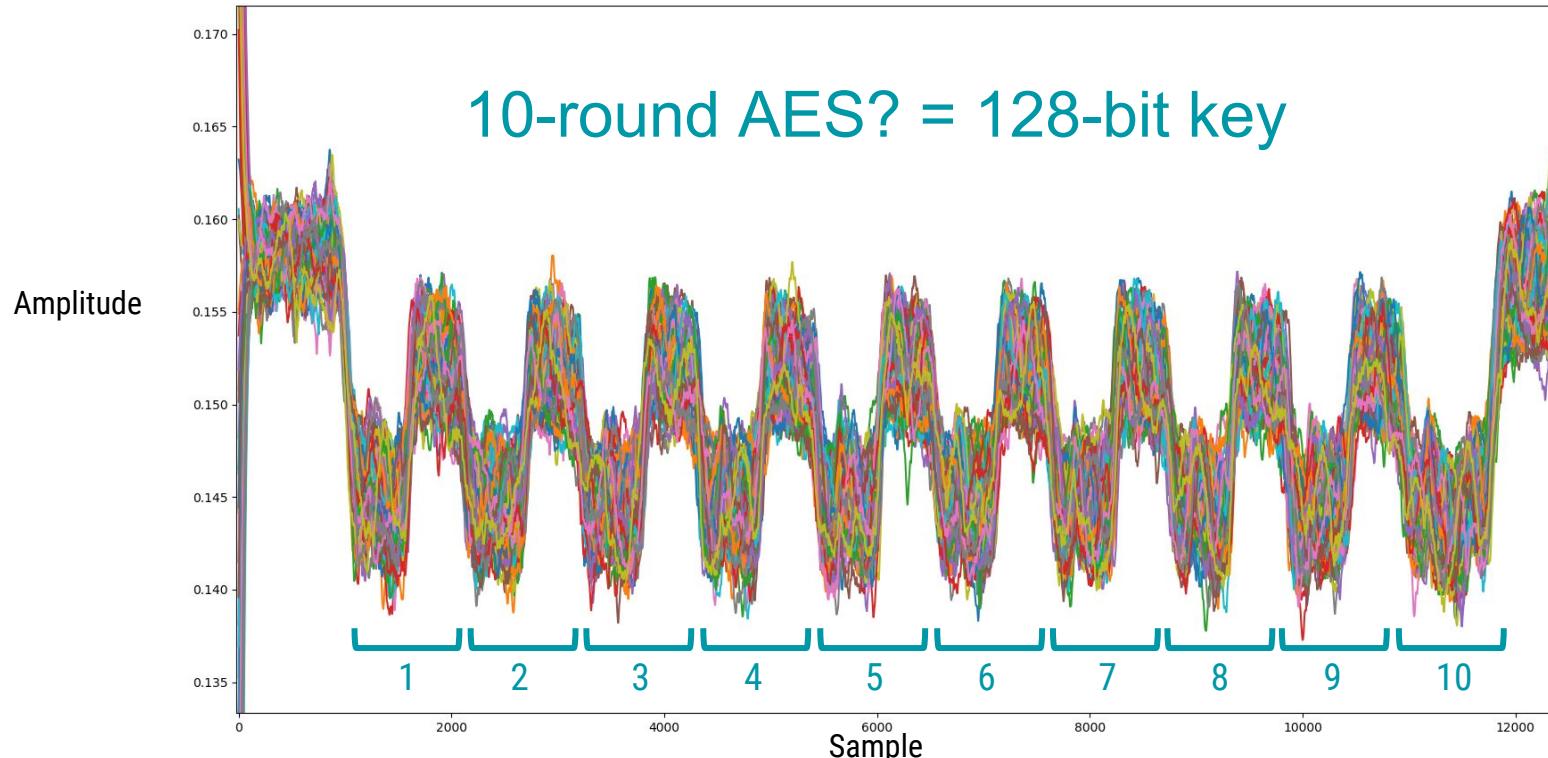
Case: AES on ATmega 328p

- After low pass filter



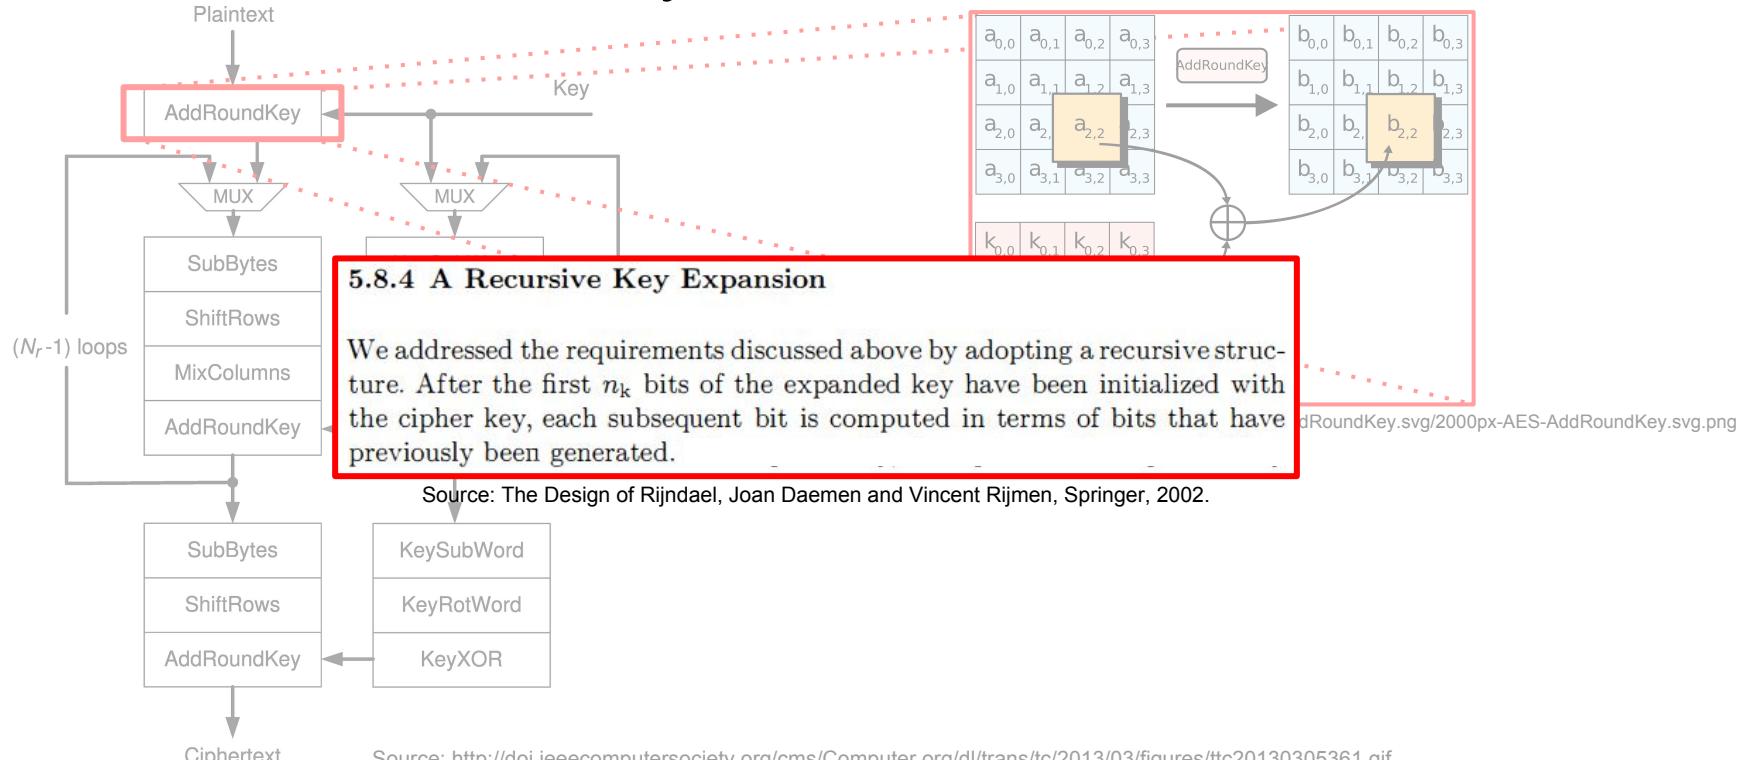
Case: AES on ATmega 328p

- After cross-correlation with reference signal



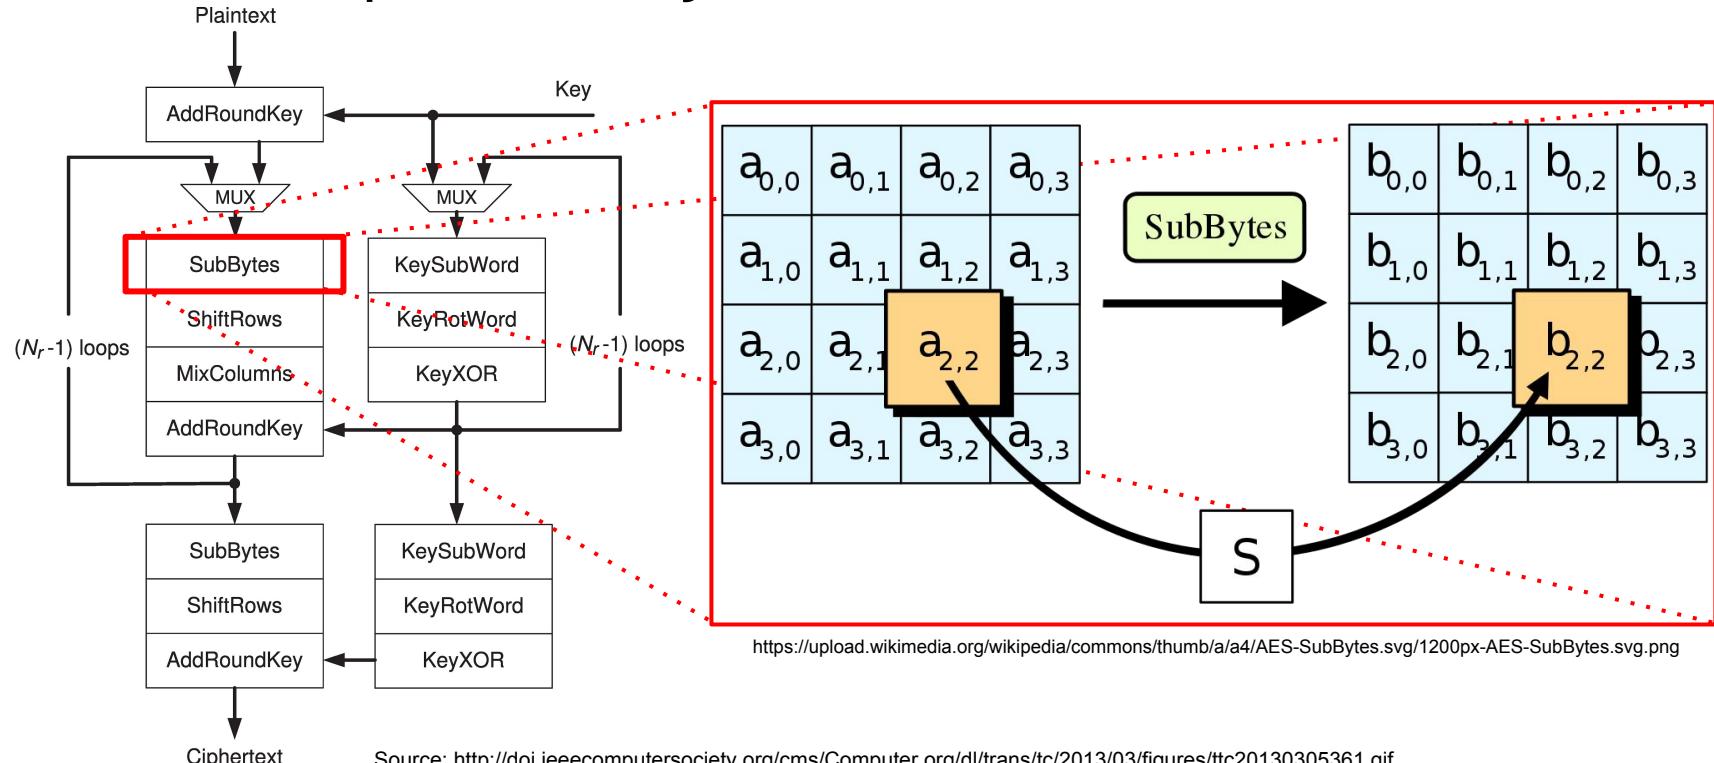
Extending our model to attack AES

- Where is the secret key in AES used?



Extending our model to attack AES

- Assume output of SubBytes is vulnerable for now



Extending our model to attack AES

- What happens inside the chip?
 - Initial state is unknown reference state R
 - After AddRoundKey and SubBytes, the state is $D = sbox[p_d \oplus k_d]$
- Current consumed ~ state changes on clock edge
 - Therefore, it's given by Hamming distance between R and D

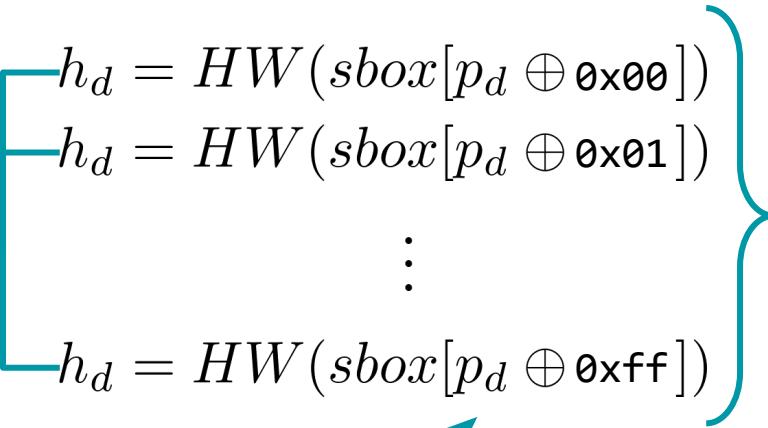
R	00100110
D	10101000

Hamming
Distance = 4

- Hamming weight also works in practice if $R = 0$

Case: AES on ATmega 328p

$$h_d = HW(sbox[p_d \oplus k_d])$$

The diagram shows a brace grouping four equations for different key byte values. The equations are:
$$h_d = HW(sbox[p_d \oplus 0x00])$$

$$h_d = HW(sbox[p_d \oplus 0x01])$$

$$\vdots$$

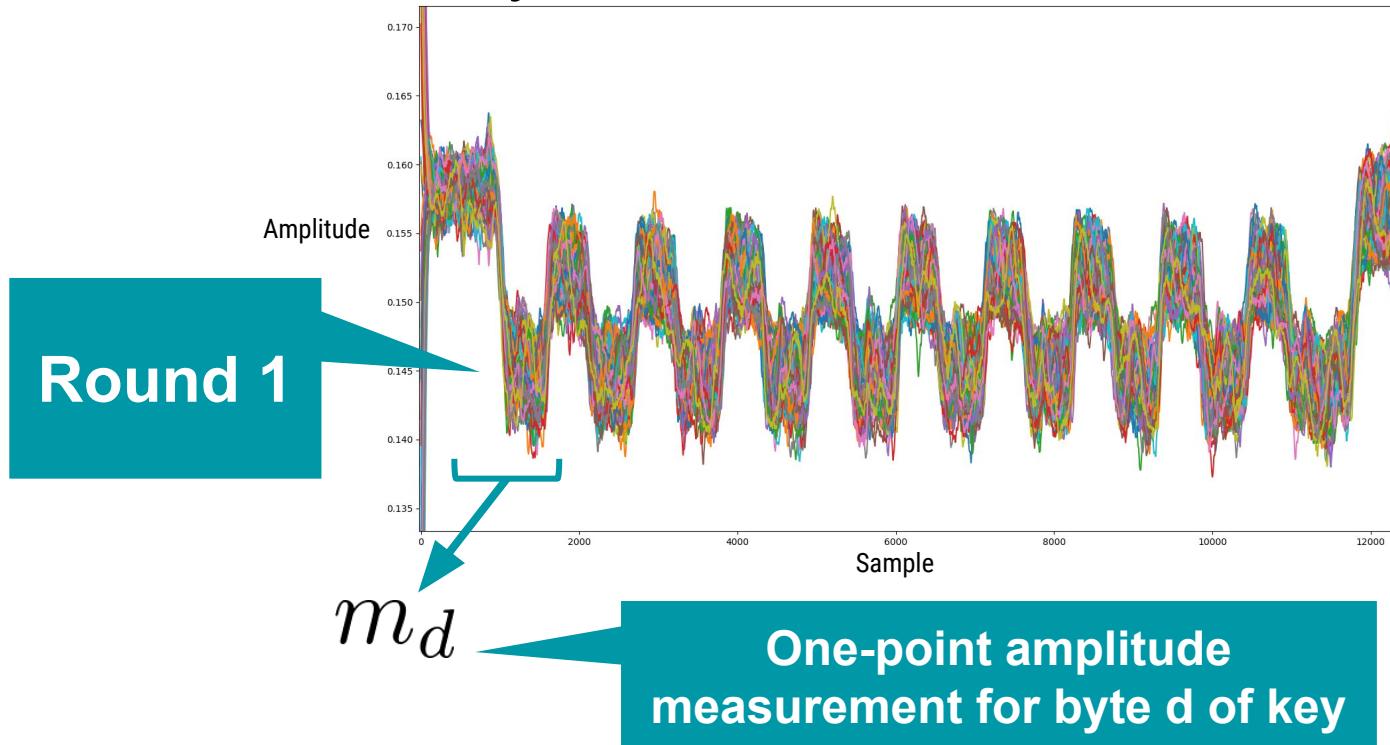
$$h_d = HW(sbox[p_d \oplus 0xff])$$

**Build models
for each
possible key
byte**

**Chosen by attacker
and varied each trace**

Case: AES on ATmega 328p

- Measure reality



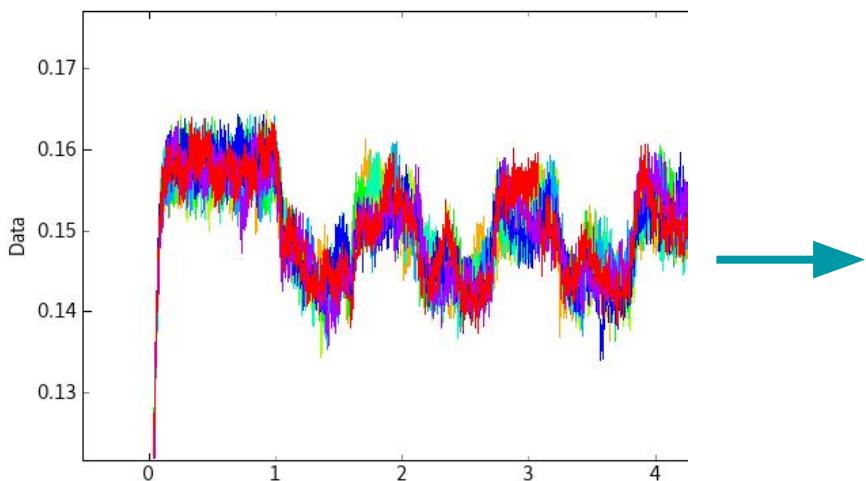
Case: AES on ATmega 328p

- Final step: correlate reality with model for each trace
- Highest correlation hypothesis is most likely key byte
- Absolute value of Pearson correlation
 - Note: only linear correlation!
- “Correlation Power Attack”

$$|\rho_{m_d, h_d}| = \left| \frac{\text{cov}(m_d, h_d)}{\sigma_{m_d} \sigma_{h_d}} \right|$$

Case: AES on ATmega 328p

- Using ChipWhisperer to perform CPA attack:



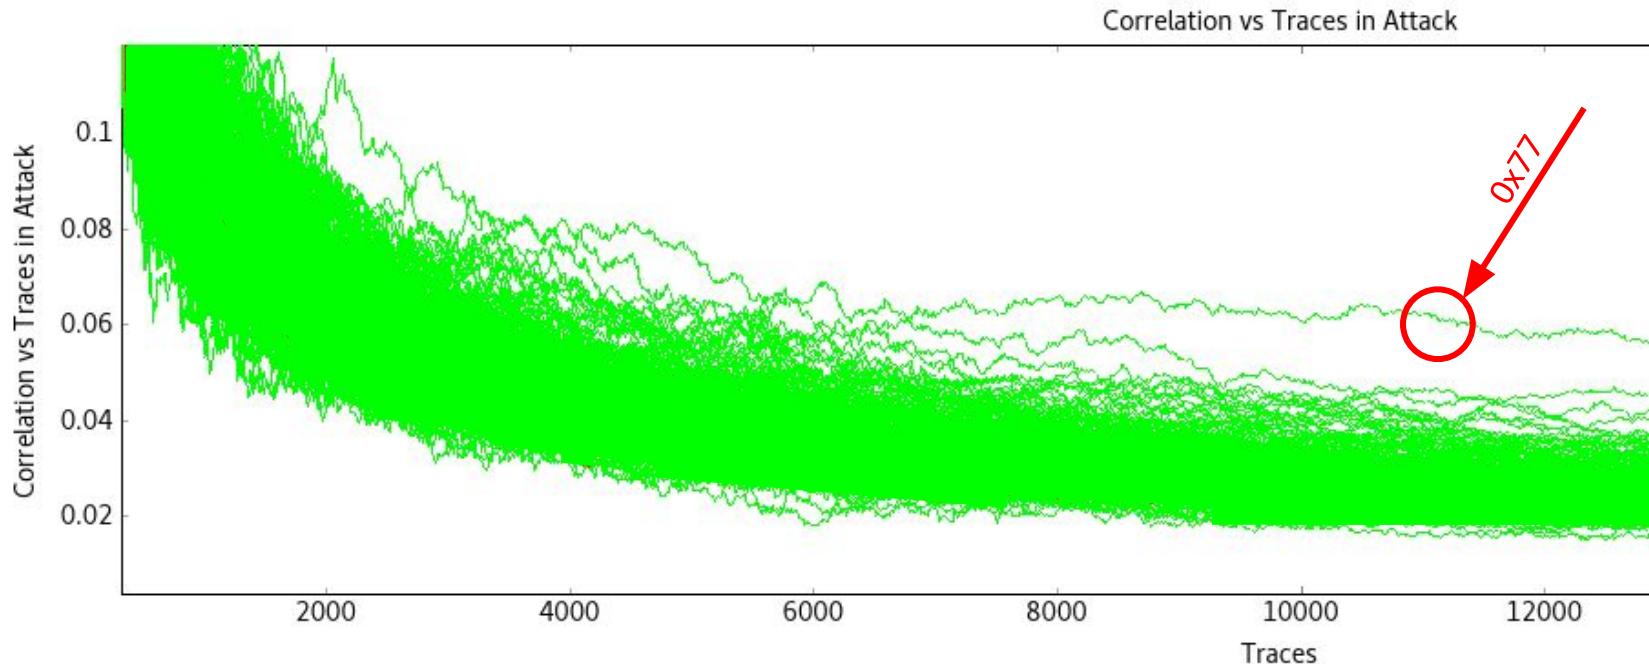
PGE	0	1	2	3	4	5	6	7
0	185 0E 0.0430	85 EB 0.0498	29 A7 0.0499	149 43 0.0563	193 00 0.0497	9 9D 0.0508	242 67 0.0599	2 D 0.04
1	C9 0.0315	15 0.0324	00 0.0449	A4 0.0372	9C 0.0322	03 0.0341	3B 0.0343	D 0.03
2	45 0.0306	3D 0.0303	69 0.0331	17 0.0325	37 0.0315	68 0.0333	40 0.0326	4 0.03
3	A3 0.0304	72 0.0292	28 0.0312	E9 0.0298	0E 0.0308	A8 0.0331	65 0.0325	D 0.03
4	19 0.0289	4B 0.0290	F3 0.0309	26 0.0298	CF 0.0293	A2 0.0329	DA 0.0323	6 0.03
5	BA 0.0279	DE 0.0288	63 0.0303	68 0.0287	3A 0.0290	92 0.0327	D7 0.0315	2 0.03
6	14 0.0278	AC 0.0287	1D 0.0301	53 0.0286	7D 0.0274	3C 0.0321	AD 0.0310	2 0.03
7	77 0.0276	87 0.0283	37 0.0299	9B 0.0283	27 0.0273	DC 0.0314	69 0.0310	4 0.03
8	BB 0.0274	89 0.0280	35 0.0299	28 0.0273	53 0.0270	47 0.0298	47 0.0308	2 0.03

Extra: SDR plugin for NewAE ChipWhisperer

Available at: http://research.edm.uhasselt.be/probysn/cw_hacky_usrp_plugin.zip

Case: AES on ATmega 328p

- Using ChipWhisperer to perform CPA attack:



Case: AES on ATmega 328p

- Using EMMA (soon-to-be open source)
 - Uses multiple cores per node and can run on multiple machines

```
Num entries: 19825
Subkey 15: elapsed: 56
Num entries: 19825

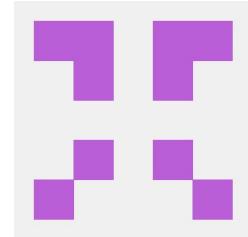
0           1           2           3           4           5           6           7           8           9           10          11          12          13          14          15
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
0.04 (0e) | 0.05 (eb) | 0.05 (a7) | 0.05 (43) | 0.04 (00) | 0.04 (9d) | 0.06 (67) | 0.04 (d2) | 0.05 (e5) | 0.05 (63) | 0.05 (cf) | 0.04 (4c) | 0.04 (5c) | 0.05 (b0) | 0.05 (77) | 0.05 (cb) |
0.02 (45) | 0.02 (15) | 0.04 (00) | 0.03 (a4) | 0.03 (37) | 0.03 (03) | 0.03 (da) | 0.03 (03) | 0.03 (32) | 0.03 (7c) | 0.03 (53) | 0.04 (30) | 0.03 (56) | 0.03 (94) | 0.03 (cc) | 0.04 (f4) |
0.02 (c9) | 0.02 (69) | 0.03 (69) | 0.03 (26) | 0.03 (ca) | 0.03 (a2) | 0.03 (3d) | 0.03 (27) | 0.02 (7d) | 0.03 (71) | 0.03 (4b) | 0.02 (d4) | 0.02 (62) | 0.03 (d7) | 0.03 (79) | 0.03 (c1) |
0.02 (77) | 0.02 (3d) | 0.03 (98) | 0.03 (9b) | 0.03 (0e) | 0.03 (3b) | 0.03 (5f) | 0.03 (d7) | 0.02 (42) | 0.03 (76) | 0.02 (3d) | 0.02 (3d) | 0.02 (1c) | 0.02 (16) | 0.03 (c3) | 0.03 (d4) |
0.02 (a0) | 0.02 (2b) | 0.03 (35) | 0.03 (53) | 0.03 (31) | 0.03 (57) | 0.03 (b9) | 0.02 (d3) | 0.02 (94) | 0.03 (dd) | 0.02 (5e) | 0.02 (a4) | 0.02 (66) | 0.02 (98) | 0.03 (eb) | 0.03 (07) |
0.02 (11) | 0.02 (01) | 0.03 (0b) | 0.02 (68) | 0.02 (80) | 0.02 (3c) | 0.02 (47) | 0.02 (4e) | 0.02 (99) | 0.03 (65) | 0.02 (7c) | 0.02 (0c) | 0.02 (7f) | 0.02 (ff) | 0.03 (4a) | 0.03 (29) |
0.02 (bb) | 0.02 (a5) | 0.03 (20) | 0.02 (a0) | 0.02 (8c) | 0.02 (6a) | 0.02 (6d) | 0.02 (2c) | 0.02 (10) | 0.03 (54) | 0.02 (b9) | 0.02 (df) | 0.02 (3f) | 0.02 (51) | 0.03 (f4) | 0.03 (8b) |
0.02 (2d) | 0.02 (08) | 0.03 (1d) | 0.02 (38) | 0.02 (c5) | 0.02 (50) | 0.02 (d7) | 0.02 (f7) | 0.02 (da) | 0.02 (24) | 0.02 (f9) | 0.02 (13) | 0.02 (9c) | 0.02 (9b) | 0.03 (de) | 0.03 (0f) |
0.02 (e2) | 0.02 (6f) | 0.03 (fb) | 0.02 (17) | 0.02 (c8) | 0.02 (ca) | 0.02 (3b) | 0.02 (48) | 0.02 (7b) | 0.02 (42) | 0.02 (70) | 0.02 (54) | 0.02 (f3) | 0.02 (22) | 0.03 (72) | 0.03 (bb) |
0.02 (41) | 0.02 (5a) | 0.02 (e4) | 0.02 (8e) | 0.02 (a3) | 0.02 (68) | 0.02 (8a) | 0.02 (7c) | 0.02 (39) | 0.02 (16) | 0.02 (25) | 0.02 (b1) | 0.02 (b5) | 0.02 (02) | 0.02 (82) | 0.03 (1d) |
0.02 (31) | 0.02 (de) | 0.02 (d0) | 0.02 (24) | 0.02 (a0) | 0.02 (92) | 0.02 (de) | 0.02 (4b) | 0.02 (c9) | 0.02 (ce) | 0.02 (be) | 0.02 (97) | 0.02 (5f) | 0.02 (6f) | 0.02 (18) | 0.03 (9c) |
0.02 (74) | 0.02 (89) | 0.02 (59) | 0.02 (dc) | 0.02 (c7) | 0.02 (20) | 0.02 (71) | 0.02 (9c) | 0.02 (a2) | 0.02 (97) | 0.02 (57) | 0.02 (4a) | 0.02 (b4) | 0.02 (74) | 0.02 (0f) | 0.03 (3b) |
0.02 (d0) | 0.02 (34) | 0.02 (e2) | 0.02 (60) | 0.02 (9c) | 0.02 (1d) | 0.02 (96) | 0.02 (3e) | 0.02 (a5) | 0.02 (de) | 0.02 (e6) | 0.02 (36) | 0.02 (45) | 0.02 (4b) | 0.02 (7d) | 0.03 (7c) |
0.02 (56) | 0.02 (19) | 0.02 (37) | 0.02 (62) | 0.02 (9f) | 0.02 (05) | 0.02 (79) | 0.02 (8e) | 0.02 (b0) | 0.02 (5c) | 0.02 (d6) | 0.02 (b2) | 0.02 (07) | 0.02 (dc) | 0.02 (20) | 0.02 (35) |
0.02 (72) | 0.02 (af) | 0.02 (0a) | 0.02 (fe) | 0.02 (11) | 0.02 (76) | 0.02 (ad) | 0.02 (b0) | 0.02 (8e) | 0.02 (ee) | 0.02 (b5) | 0.02 (b6) | 0.02 (8d) | 0.02 (41) | 0.02 (b2) | 0.02 (09) |
0.02 (ce) | 0.02 (75) | 0.02 (a3) | 0.02 (02) | 0.02 (d7) | 0.02 (78) | 0.02 (32) | 0.02 (75) | 0.02 (eb) | 0.02 (08) | 0.02 (17) | 0.02 (6f) | 0.02 (c8) | 0.02 (ab) | 0.02 (21) | 0.02 (5c) |
0.02 (e3) | 0.02 (f3) | 0.02 (ba) | 0.02 (10) | 0.02 (1e) | 0.02 (cc) | 0.02 (e9) | 0.02 (9d) | 0.02 (5d) | 0.02 (19) | 0.02 (b2) | 0.02 (82) | 0.02 (a4) | 0.02 (d3) | 0.02 (31) | 0.02 (3c) |
0.02 (d4) | 0.02 (05) | 0.02 (3f) | 0.02 (4a) | 0.02 (17) | 0.02 (ce) | 0.02 (f3) | 0.02 (f6) | 0.02 (40) | 0.02 (01) | 0.02 (5f) | 0.02 (e1) | 0.02 (29) | 0.02 (fb) | 0.02 (d2) | 0.02 (69) |
0.02 (3b) | 0.02 (e5) | 0.02 (af) | 0.02 (bb) | 0.02 (20) | 0.02 (c7) | 0.02 (35) | 0.02 (72) | 0.02 (d9) | 0.02 (34) | 0.02 (fe) | 0.02 (0e) | 0.02 (f7) | 0.02 (e3) | 0.02 (60) | 0.02 (0a) |
0.02 (03) | 0.02 (ff) | 0.02 (c0) | 0.02 (77) | 0.02 (b1) | 0.02 (38) | 0.02 (85) | 0.02 (36) | 0.02 (79) | 0.02 (84) | 0.02 (20) | 0.02 (27) | 0.02 (96) | 0.02 (c8) | 0.02 (05) | 0.02 (d1) |
0e eb a7 43 00 9d 67 d2 e5 63 cf 4c 5c b0 77 cb
Cleaning up
[probyns@compute-4 emma]$
```

Closing statements



- All my finished research is open source

Decoder: <https://github.com/rpp0/gr-lora>



Fingerprinting: <https://github.com/rpp0/lora-phy-fingerprinting>

ChipWhisperer plugin: http://research.edm.uhasselt.be/probyns/cw_hacky_usrp_plugin.zip

- Some of my current research directions
 - Relation to machine learning → loss function and features vs. correlation
 - Can we improve the state of the art in this way?
 - Increasing the range of EM attacks
 - Analyzing below the noise floor, custom antenna designs, etc.
 - Open to collaborations!

Further reading

- Here are some related papers which I found interesting

Fingerprinting

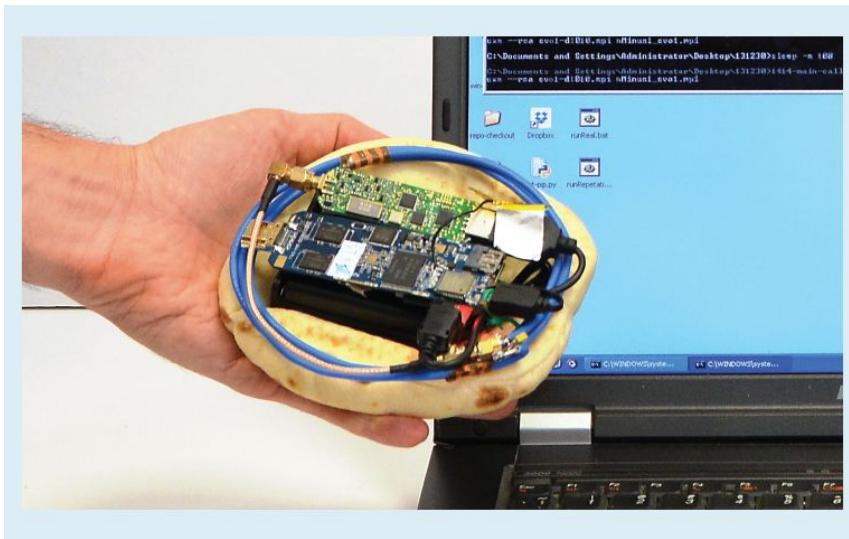
- *Why MAC address randomization is not enough...* (Mathy Vanhoef et al.)
- *Challenges to PHY anonymity for Wi-Fi* (Peter Iannucci)
- *Convolutional Radio Modulation Recognition...* (Timothy O'Shea et al.)
- *Unsupervised Learning on Neural Network Outputs* (Yao Lu et al.)
- *Device Fingerprinting in Wireless Networks...* (Qiang Xu et al.)

EM side-channel attacks

- *Correlation Power Analysis with a Leakage Model* (Eric Brier et al.)
- *Enhancing Electromagnetic Side-Channel Analysis in...* (David P. Montminy.)
- *NewAE Wiki page* (https://wiki.newae.com/Main_Page)
- *Power Analysis Attacks against IEEE 802.15.4 Nodes* (Colin O'Flynn et al.)

Other nice examples of EM side channel attacks

Fully extract decryption keys, by measuring the laptop's chassis potential during decryption of a chosen ciphertext.

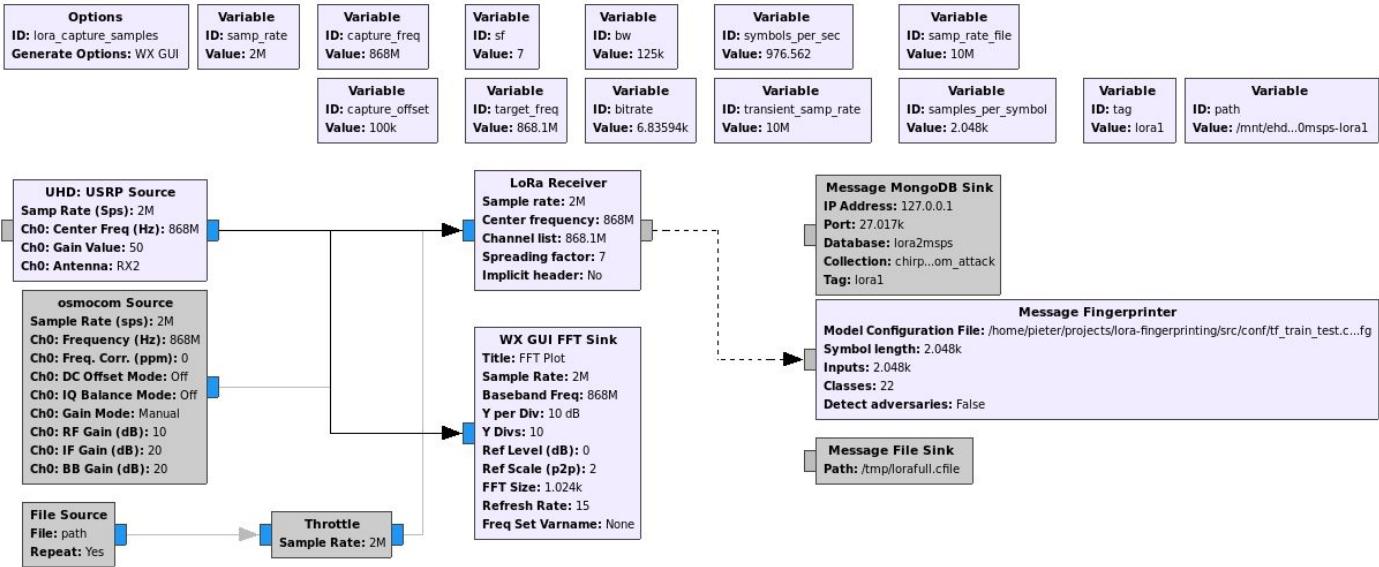


Full extraction of ECDSA secret signing keys from OpenSSL and CoreBitcoin running on iOS devices.



Source: <https://www.tau.ac.il/~tromer/handoff/>

Demo



Questions?

pieter.robyns@uhasselt.be



UHASSELT

EDM

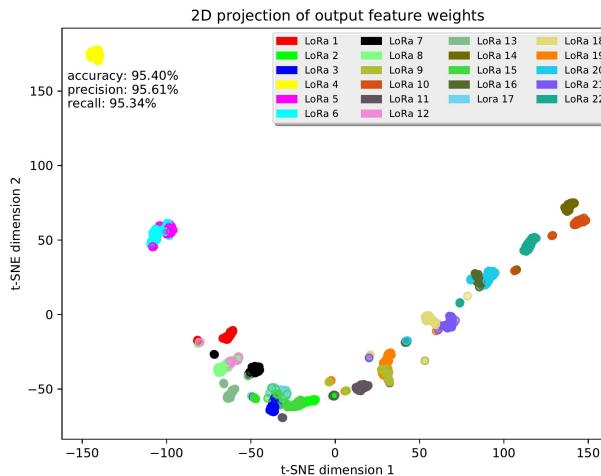


Research Foundation
Flanders
Opening new horizons

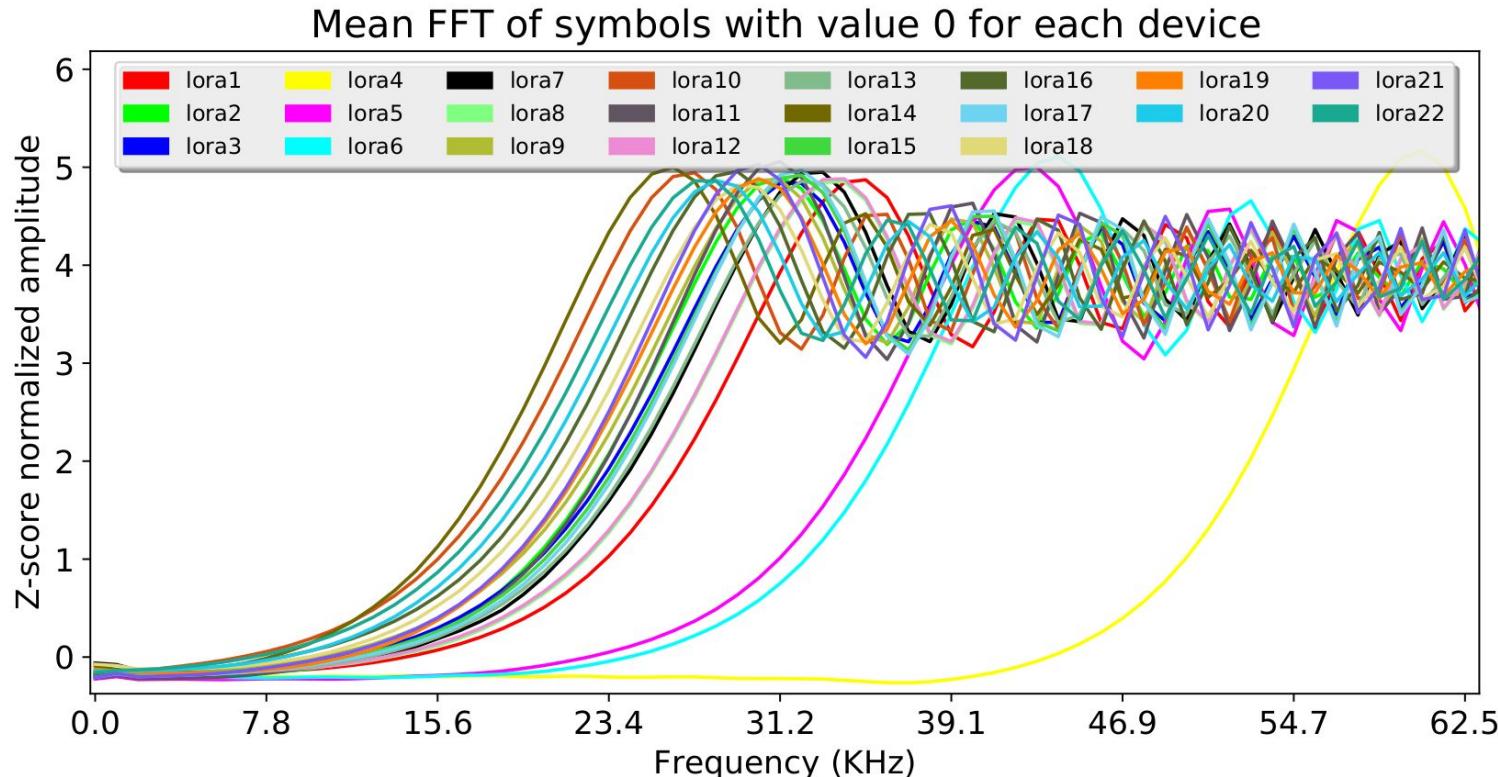
Extra slides

But wait, what about devices that we can't train?

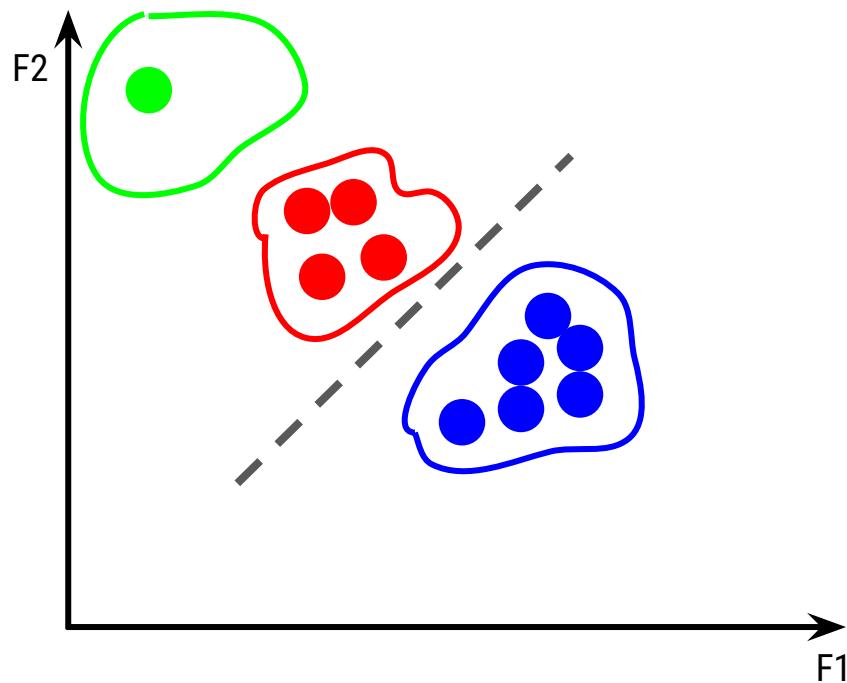
- Technique called zero shot classification
 - Learn “attributes” during training
 - Describe unseen devices using learned attributes
 - Example: cluster on neural network outputs that was trained with a number known LoRa devices



But wait, what about devices that we can't train?

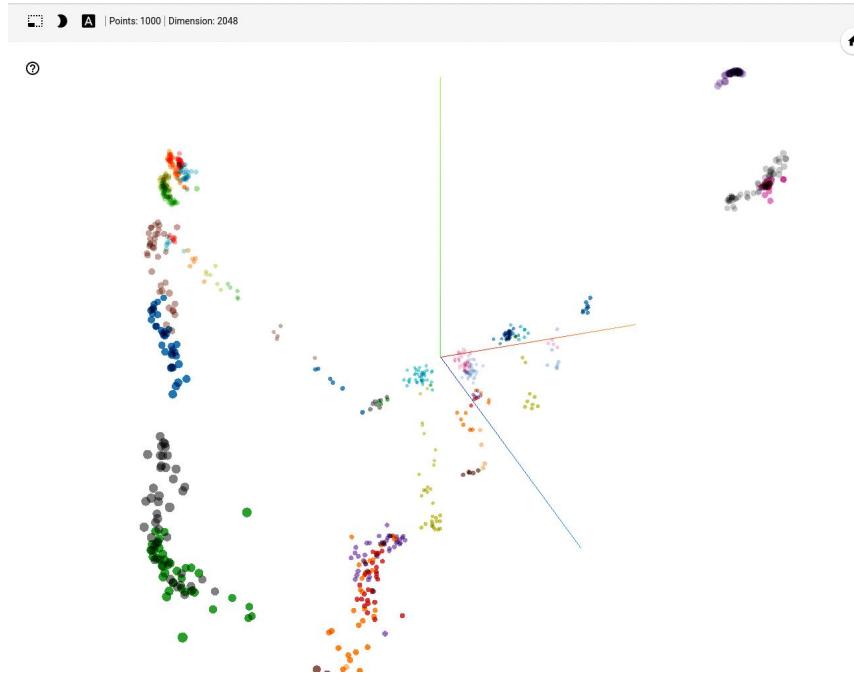


But wait, what about devices that we can't train?

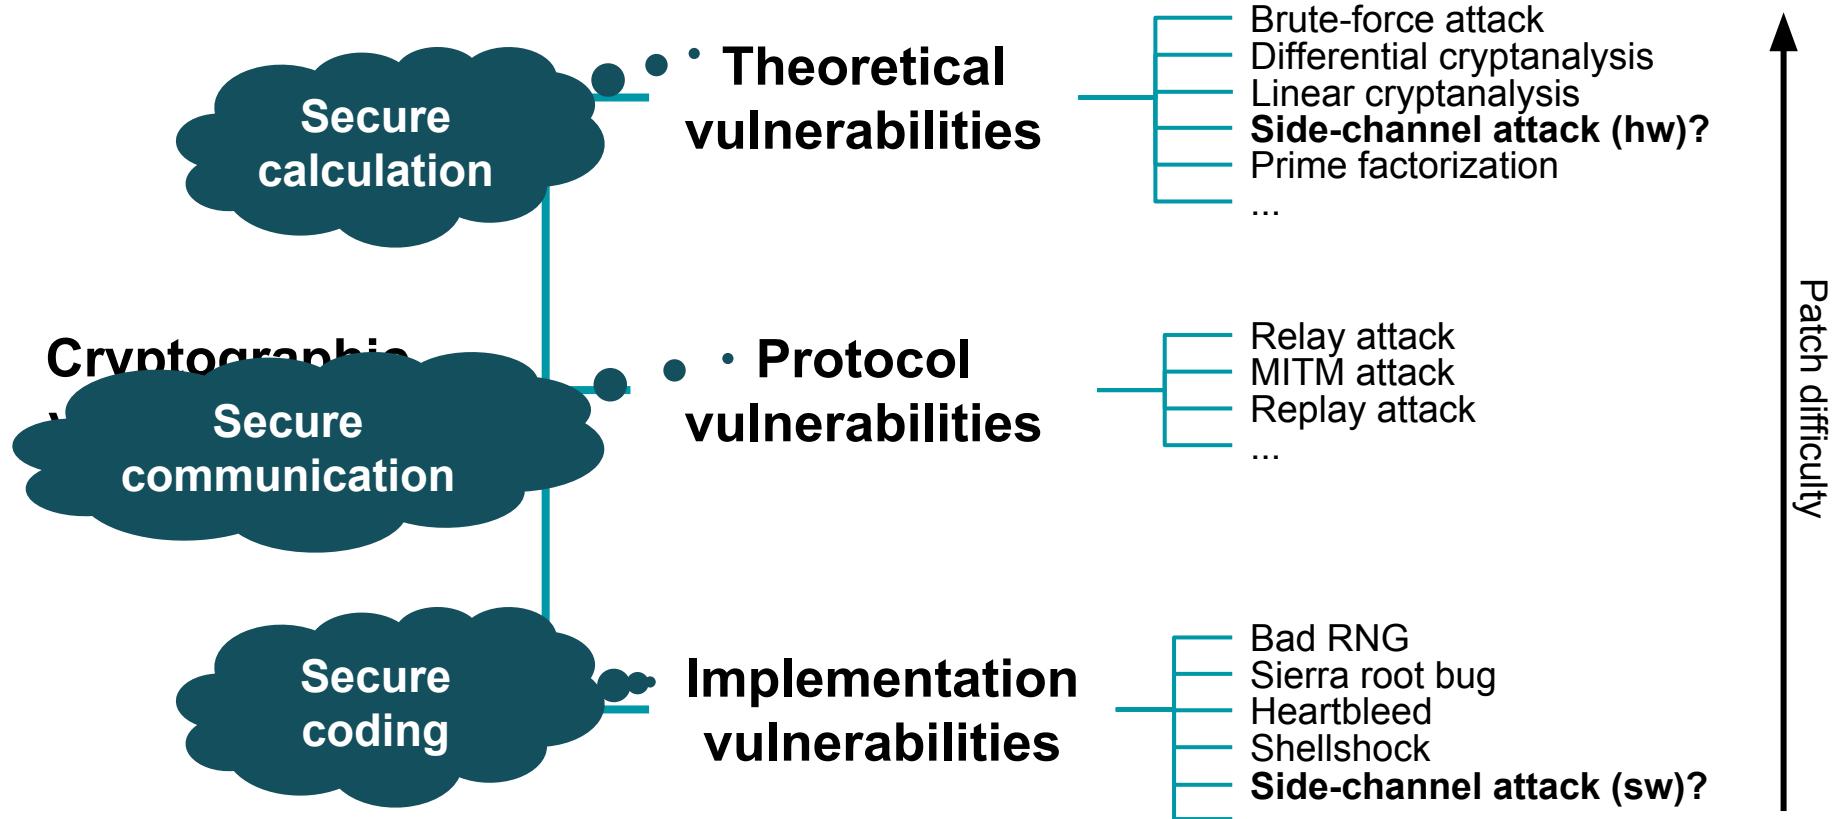


Visualizing the raw data

- Visualizing the signal using Principal Component Analysis (PCA):



SCAs within the vulnerability landscape



SCAs within the vulnerability landscape



Theoretical vulnerabilities

- Brute-force attack
- Differential cryptanalysis
- Linear cryptanalysis
- Side-channel attack (hw)?**
- Prime factorization
- ...

Should the hardware or theoretical design automatically mitigate dangerous calculations (temperature, radiation,...) or should the programmer implement the theoretical design in such a way that exploitation is not possible?



Implementation vulnerabilities

- Bad RNG
- Sierra root bug
- Heartbleed
- Shellshock
- Side-channel attack (sw)?**
- ...